

# deep-cnn-image-classifier

February 16, 2024

## 1 Deep CNN Image Classifier with ANY Images.

### 1.0.1 1. Install Dependencies and Setup

```
[1]: import tensorflow as tf
import os

# Avoid OOM errors by setting GPU Memory Consumption Growth
gpus = tf.config.experimental.list_physical_devices('GPU')
for gpu in gpus:
    tf.config.experimental.set_memory_growth(gpu, True)
```

```
[2]: gpus
```

```
[2]: [PhysicalDevice(name='/physical_device:GPU:0', device_type='GPU')]
```

```
[3]: gpus = tf.config.experimental.list_physical_devices('CPU')
gpus
```

```
[3]: [PhysicalDevice(name='/physical_device:CPU:0', device_type='CPU')]
```

### 1.1 2. Remove Dodgy images

```
[4]: import cv2
import imghdr
```

```
[5]: data_dir = 'drive/MyDrive/Project 13 Classification/data'
```

```
[6]: # cheking no of images in the dataset
total_images = 0

# Loop over each subdirectory
for folder_name in os.listdir(data_dir):
    folder_path = os.path.join(data_dir, folder_name)

    # Ensure the path is a directory before counting images
    if os.path.isdir(folder_path):
```

```

        num_images = len(os.listdir(folder_path))
        total_images += num_images

print(f'Total number of images: {total_images}')

```

Total number of images: 225

```
[7]: image_exts = ['jpeg', 'jpg', 'bmp', 'png']
```

This code is used to read and process images from a directory. It checks the file type of each image and deletes any images that are not of the expected file type or cannot be read.

```
[8]: for image_class in os.listdir(data_dir):
    for image in os.listdir(os.path.join(data_dir, image_class)):
        image_path = os.path.join(data_dir, image_class, image)
        try:
            img = cv2.imread(image_path)
            tip = imghdr.what(image_path)
            if tip not in image_exts:
                print('Image not ext list {}'.format(image_path))
                os.remove(image_path)
        except Exception as e:
            print('Issue with image {}'.format(image_path))
            os.remove(image_path)

```

### 1.1.1 3. Load Data

```
[9]: import numpy as np
    from matplotlib import pyplot as plt

```

```
[10]: data = 'drive/MyDrive/Project 13 Classification/data'
```

Using the Keras.Utilit function to preprocess the data into shape and size that I need in its default format

```
[11]: data = tf.keras.utils.image_dataset_from_directory(data)
```


Found 225 files belonging to 2 classes.

```
[12]: data_iterator = data.as_numpy_iterator()
```

```
[13]: batch = data_iterator.next()
```

```
[14]: fig, ax = plt.subplots(ncols = 4, figsize = (20,20))
    for idx, img in enumerate(batch[0][:4]):
        ax[idx].imshow(img.astype(int))
        ax[idx].title.set_text(batch[1][idx])

```



1 represents Sad people and 0 represents Happy people

#### 1.1.2 4. Scale The data

```
[15]: data = data.map(lambda x, y: (x/255, y))
```

```
[16]: data.as_numpy_iterator().next()
```

```
[16]: (array([[[[0.08583793, 0.12127757, 0.23587623],
               [0.0723269 , 0.11285233, 0.25447303],
               [0.05271906, 0.10660999, 0.28235295],
               ...,
               [0.1279565 , 0.1279565 , 0.2769761 ],
               [0.12941177, 0.12941177, 0.2714614 ],
               [0.12941177, 0.12941177, 0.27058825]],

               [[0.08583793, 0.11662266, 0.22897713],
               [0.07282658, 0.10387561, 0.24325214],
               [0.06012222, 0.09804967, 0.27029926],
               ...,
               [0.12612872, 0.12571232, 0.27556473],
               [0.12716758, 0.12716758, 0.26921722],
               [0.12716758, 0.12716758, 0.26834404]],

               [[0.09159287, 0.115437 , 0.22269455],
               [0.08308484, 0.10066916, 0.23348859],
               [0.08546917, 0.09615503, 0.25874773],
               ...,
               [0.12497147, 0.12220091, 0.27359277],
               [0.1254902 , 0.12323835, 0.26841506],
               [0.1254902 , 0.12323835, 0.26779258]],

               ...,

               [[0.1102405 , 0.13376991, 0.2749464 ],
               [0.10590253, 0.12899536, 0.27104497],
```

```

[0.10048043, 0.12008827, 0.2691079 ],
...,
[0.13012198, 0.15602367, 0.3416973 ],
[0.12378217, 0.15092191, 0.33193654],
[0.12000613, 0.14745711, 0.32784927]],

[[0.09741212, 0.11926415, 0.26379538],
[0.09500138, 0.11660357, 0.2616345 ],
[0.09259064, 0.11219849, 0.26121807],
...,
[0.14298128, 0.17069827, 0.35683978],
[0.13612132, 0.16499986, 0.34607843],
[0.13234529, 0.16147365, 0.3418658 ]],

[[0.08627451, 0.10588235, 0.25490198],
[0.08627451, 0.10588235, 0.25490198],
[0.08627451, 0.10588235, 0.25490198],
...,
[0.15237439, 0.18374693, 0.36806065],
[0.14509805, 0.1764706 , 0.35729933],
[0.141322 , 0.17269455, 0.3530867 ]]],

[[[0.1799326 , 0.7448836 , 0.7682292 ],
[0.1799326 , 0.7448836 , 0.7682292 ],
[0.1799326 , 0.7448836 , 0.7682292 ],
...,
[0.1799326 , 0.7448836 , 0.7682292 ],
[0.1799326 , 0.7448836 , 0.7682292 ],
[0.1799326 , 0.7448836 , 0.7682292 ]],

[[0.05882353, 0.7058824 , 0.73333335],
[0.05882353, 0.7058824 , 0.73333335],
[0.05882353, 0.7058824 , 0.73333335],
...,
[0.05882353, 0.7058824 , 0.73333335],
[0.05882353, 0.7058824 , 0.73333335],
[0.05882353, 0.7058824 , 0.73333335]],

[[0.05882353, 0.7058824 , 0.73333335],
[0.05882353, 0.7058824 , 0.73333335],
[0.05882353, 0.7058824 , 0.73333335],
...,
[0.05882353, 0.7058824 , 0.73333335],
[0.05882353, 0.7058824 , 0.73333335],
[0.05882353, 0.7058824 , 0.73333335]],

```

```

...,

[[0.05882353, 0.7058824 , 0.73333335],
 [0.05882353, 0.7058824 , 0.73333335],
 [0.05882353, 0.7058824 , 0.73333335],
 ...,
 [0.05882353, 0.7058824 , 0.73333335],
 [0.05882353, 0.7058824 , 0.73333335],
 [0.05882353, 0.7058824 , 0.73333335]],

[[0.05882353, 0.7058824 , 0.73333335],
 [0.05882353, 0.7058824 , 0.73333335],
 [0.05882353, 0.7058824 , 0.73333335],
 ...,
 [0.05882353, 0.7058824 , 0.73333335],
 [0.05882353, 0.7058824 , 0.73333335],
 [0.05882353, 0.7058824 , 0.73333335]],

[[0.05882353, 0.7058824 , 0.73333335],
 [0.05882353, 0.7058824 , 0.73333335],
 [0.05882353, 0.7058824 , 0.73333335],
 ...,
 [0.05882353, 0.7058824 , 0.73333335],
 [0.05882353, 0.7058824 , 0.73333335],
 [0.05882353, 0.7058824 , 0.73333335]]],

[[[0.9421875 , 0.9186581 , 0.9244792 ],
 [0.94509804, 0.9254902 , 0.9106924 ],
 [0.93838847, 0.9187806 , 0.8952512 ],
 ...,
 [0.9499344 , 0.93424815, 0.92248344],
 [0.9448065 , 0.92912024, 0.91735554],
 [0.94509804, 0.92941177, 0.91764706]],

[[0.9421875 , 0.9186581 , 0.9244792 ],
 [0.94509804, 0.9254902 , 0.9106924 ],
 [0.93838847, 0.9187806 , 0.8952512 ],
 ...,
 [0.94231004, 0.92662376, 0.91485906],
 [0.9499081 , 0.9342218 , 0.9224571 ],
 [0.94509804, 0.92941177, 0.91764706]],

[[0.9421875 , 0.9186581 , 0.9244792 ],
 [0.94509804, 0.9254902 , 0.9106924 ],
 [0.93838847, 0.9187806 , 0.8952512 ],
 ...,

```

```

[0.94650733, 0.93082106, 0.91905636],
[0.9476103 , 0.93192405, 0.92015934],
[0.94509804, 0.92941177, 0.91764706]],

...,

[[[0.11997549, 0.15919118, 0.16703431],
  [0.11352395, 0.15273964, 0.16058278],
  [0.11764706, 0.14901961, 0.16078432],
  ...,
  [0.4392157 , 0.47058824, 0.6179534 ],
  [0.43836263, 0.4703867 , 0.6008885 ],
  [0.43308824, 0.4762255 , 0.6017157 ]],

  [[0.11368001, 0.1528957 , 0.16073884],
   [0.1106924 , 0.1499081 , 0.15775123],
   [0.11378676, 0.14515932, 0.15692402],
   ...,
   [0.42371324, 0.46169052, 0.60631126],
   [0.42454043, 0.45983455, 0.58924633],
   [0.4272059 , 0.47034314, 0.59583336]],

  [[0.11014093, 0.14935662, 0.15719976],
   [0.12484585, 0.16406155, 0.17190468],
   [0.11335354, 0.14472608, 0.15649079],
   ...,
   [0.4208946 , 0.45897672, 0.60355395],
   [0.41786152, 0.46099877, 0.58648896],
   [0.4143995 , 0.45753676, 0.58302695]]],

...,

[[[0.02331495, 0.01939338, 0.00370711],
  [0.04571078, 0.03002451, 0.0182598 ],
  [0.05892549, 0.04716079, 0.02755295],
  ...,
  [0.04313726, 0.06666667, 0.06666667],
  [0.04283711, 0.05852338, 0.07028808],
  [0.03529412, 0.05251225, 0.06427696]],

  [[0.05690583, 0.04121955, 0.02945485],
   [0.05223652, 0.03655025, 0.02478554],
   [0.05900735, 0.04724265, 0.0276348 ],
   ...,
   [0.04313726, 0.06666667, 0.06666667],

```

```

[0.04683239, 0.06251867, 0.07428338],
[0.03534917, 0.05256731, 0.06433202]],

[[0.0836397 , 0.06403186, 0.05226716],
[0.04712058, 0.0314343 , 0.0196696 ],
[0.04808948, 0.04416791, 0.02456007],
...,
[0.04024634, 0.06377576, 0.06377576],
[0.04705882, 0.0627451 , 0.07450981],
[0.03953211, 0.05675025, 0.06851496]],

...,

[[0.04947917, 0.0573223 , 0.03771446],
[0.04963235, 0.05747549, 0.04436275],
[0.0579767 , 0.06581984, 0.05405513],
...,
[0.10283826, 0.10283826, 0.07118136],
[0.06525687, 0.07702158, 0.0495706 ],
[0.06935269, 0.0811174 , 0.05366642]],

[[0.05775123, 0.0616728 , 0.04053309],
[0.05165441, 0.05557598, 0.03596814],
[0.05508579, 0.05508579, 0.04724265],
...,
[0.07337622, 0.06945466, 0.04984681],
[0.05165441, 0.07126226, 0.04773284],
[0.05729167, 0.07689951, 0.05183823]],

[[0.06670544, 0.07062701, 0.04403387],
[0.05165441, 0.05557598, 0.03204657],
[0.05508579, 0.05900735, 0.03939951],
...,
[0.06481837, 0.06873995, 0.0491321 ],
[0.05098039, 0.07058824, 0.04705882],
[0.0562275 , 0.07583535, 0.05230593]]],

[[[0.0870481 , 0.13043046, 0.0755285 ],
[0.0508004 , 0.09785922, 0.04295726],
[0.07494638, 0.12984835, 0.07102481],
...,
[0.887343 , 0.89101946, 0.92116654],
[0.8901961 , 0.89411765, 0.91158473],
[0.89411765, 0.8901961 , 0.9135608 ]],

[[0.10459942, 0.14828432, 0.08144148],

```

```

[0.08870251, 0.1396829 , 0.07295496],
[0.09021523, 0.1451172 , 0.08620941],
...,
[0.8877643 , 0.8940985 , 0.91099113],
[0.89025736, 0.89411765, 0.9019608 ],
[0.89411765, 0.89411765, 0.9019608 ]],

[[[0.13482307, 0.18221508, 0.10087699],
[0.12175628, 0.17273667, 0.09932981],
[0.09508655, 0.15095742, 0.0839231 ]],
...,
[0.8877451 , 0.89166665, 0.90832186],
[0.89270836, 0.8936466 , 0.9059858 ],
[0.89270836, 0.89411765, 0.90477943]]],

...,

[[[0.07266391, 0.12072993, 0.08167126],
[0.09770986, 0.13974418, 0.10304075],
[0.06666284, 0.10220205, 0.06813343],
...,
[0.2565602 , 0.30607 , 0.21317784],
[0.25484067, 0.2893995 , 0.18572304],
[0.22944623, 0.26973805, 0.15971585]]],

[[[0.05924479, 0.0868796 , 0.05544577],
[0.06571309, 0.09334788, 0.06191406],
[0.0639208 , 0.09511719, 0.06249617],
...,
[0.2859107 , 0.33419502, 0.2503715 ],
[0.1971852 , 0.22860754, 0.13434054],
[0.20478325, 0.24375384, 0.13787147]]],

[[[0.0313151 , 0.04851026, 0.02105928],
[0.02414982, 0.04375766, 0.01630668],
[0.05180377, 0.08293122, 0.05033318],
...,
[0.27328813, 0.32157245, 0.23951823],
[0.21140854, 0.2427811 , 0.14939874],
[0.17912453, 0.21785003, 0.12005591]]],

[[[0.22218138, 0.22218138, 0.19007353],
[0.26813725, 0.23284313, 0.2125 ],
[0.3348039 , 0.25674018, 0.29093137],
...,
[0.41789216, 0.40294117, 0.45747548],

```



```

[0.39191177, 0.34485295, 0.3997549 ],
[0.41066176, 0.36360294, 0.37144607]],

[[0.24093138, 0.21960784, 0.21629901],
[0.2617647 , 0.22254902, 0.22316177],
[0.31875 , 0.25490198, 0.29693627],
...,
[0.4302696 , 0.41458333, 0.46948528],
[0.41507354, 0.36409312, 0.43075982],
[0.4509804 , 0.40563726, 0.43590686]],

[[0.22058824, 0.19105393, 0.22683823],
[0.25808823, 0.21495098, 0.24264705],
[0.3262255 , 0.28125 , 0.34031862],
...,
[0.39791667, 0.37438726, 0.42928922],
[0.41703433, 0.37389705, 0.44056374],
[0.44154412, 0.39240196, 0.45355392]],

...,

[[0.5686275 , 0.49019608, 0.53333336],
[0.53933823, 0.4569853 , 0.4922794 ],
[0.56421566, 0.4897059 , 0.5132353 ],
...,
[0.41335785, 0.32316175, 0.35502452],
[0.41458333, 0.32230392, 0.3677696 ],
[0.4209559 , 0.34068626, 0.3822304 ]],

[[0.5283088 , 0.4627451 , 0.50306374],
[0.5598039 , 0.47855392, 0.5160539 ],
[0.54791665, 0.47340685, 0.49693626],
...,
[0.38308823, 0.29571077, 0.3107843 ],
[0.40551472, 0.3052696 , 0.32769608],
[0.40563726, 0.31323528, 0.3345588 ]],

[[0.50980395, 0.44705883, 0.49019608],
[0.5254902 , 0.45882353, 0.49803922],
[0.56531864, 0.49865195, 0.5221814 ],
...,
[0.3925245 , 0.30232844, 0.33370098],
[0.41066176, 0.2930147 , 0.32438725],
[0.40355393, 0.29767156, 0.32512254]]], dtype=float32),
array([0, 0, 1, 1, 0, 1, 0, 0, 0, 1, 0, 1, 0, 1, 0, 1, 1, 0, 1, 1, 0, 1,
0, 0, 1, 0, 1, 0, 0, 0, 0, 1], dtype=int32))

```

### 1.1.3 5. Split Data

```
[17]: train_size = int(len(data)*.7)
      val_size = int(len(data)*.2)
      test_size = int(len(data)*.1)
```

```
[18]: train_size
```

```
[18]: 5
```

```
[19]: train = data.take(train_size)
      val = data.skip(train_size).take(val_size)
      test = data.skip(train_size + val_size).take(test_size)
```

### 1.2 6. Build Deep Learning Model

```
[20]: train
```

```
[20]: <_TakeDataset element_spec=(TensorSpec(shape=(None, 256, 256, 3),
      dtype=tf.float32, name=None), TensorSpec(shape=(None,), dtype=tf.int32,
      name=None))>
```

```
[21]: from tensorflow.keras.models import Sequential
      from tensorflow.keras.layers import Conv2D, MaxPooling2D, Dense, Flatten, Dropout
```

```
[22]: model = Sequential()
```

```
[23]: model.add(Conv2D(16, (3,3), 1, activation='relu', input_shape=(256,256,3)))
      model.add(MaxPooling2D())
      model.add(Conv2D(32, (3,3), 1, activation='relu'))
      model.add(MaxPooling2D())
      model.add(Conv2D(16, (3,3), 1, activation='relu'))
      model.add(MaxPooling2D())
      model.add(Flatten())
      model.add(Dense(256, activation='relu'))
      model.add(Dense(1, activation='sigmoid'))
```

```
[24]: # Compiling the model
      model.compile('adam', loss=tf.losses.BinaryCrossentropy(), metrics=['accuracy'])
```

```
[25]: # Models summary
      model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
--------------	--------------	---------

```

=====
conv2d (Conv2D)                (None, 254, 254, 16)    448

max_pooling2d (MaxPooling2D)   (None, 127, 127, 16)    0

conv2d_1 (Conv2D)              (None, 125, 125, 32)    4640

max_pooling2d_1 (MaxPooling2D) (None, 62, 62, 32)      0

conv2d_2 (Conv2D)              (None, 60, 60, 16)      4624

max_pooling2d_2 (MaxPooling2D) (None, 30, 30, 16)      0

flatten (Flatten)              (None, 14400)           0

dense (Dense)                  (None, 256)              3686656

dense_1 (Dense)                (None, 1)                257

=====
Total params: 3696625 (14.10 MB)
Trainable params: 3696625 (14.10 MB)
Non-trainable params: 0 (0.00 Byte)
-----

```

### 1.3 7. Train

```
[26]: #save the view the logs
logdir = 'logs'
```

```
[27]: tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir = logdir)
```

```
[28]: hist = model.fit(train, epochs = 60, validation_data = val, callbacks = [
    ↪ tensorboard_callback])
```

```

Epoch 1/60
5/5 [=====] - 12s 877ms/step - loss: 1.3285 - accuracy:
0.5375 - val_loss: 1.1755 - val_accuracy: 0.3750
Epoch 2/60
5/5 [=====] - 6s 820ms/step - loss: 0.8052 - accuracy:
0.5375 - val_loss: 0.6359 - val_accuracy: 0.6562
Epoch 3/60
5/5 [=====] - 8s 850ms/step - loss: 0.7123 - accuracy:
0.5625 - val_loss: 0.6619 - val_accuracy: 0.5938
Epoch 4/60

```

5/5 [=====] - 8s 1s/step - loss: 0.6366 - accuracy:  
 0.6187 - val\_loss: 0.5146 - val\_accuracy: 0.7812  
 Epoch 5/60  
 5/5 [=====] - 7s 840ms/step - loss: 0.6036 - accuracy:  
 0.5875 - val\_loss: 0.5359 - val\_accuracy: 0.6875  
 Epoch 6/60  
 5/5 [=====] - 9s 1s/step - loss: 0.5356 - accuracy:  
 0.7125 - val\_loss: 0.5782 - val\_accuracy: 0.5938  
 Epoch 7/60  
 5/5 [=====] - 6s 826ms/step - loss: 0.4605 - accuracy:  
 0.7688 - val\_loss: 0.3986 - val\_accuracy: 0.7500  
 Epoch 8/60  
 5/5 [=====] - 10s 1s/step - loss: 0.3928 - accuracy:  
 0.8313 - val\_loss: 0.3784 - val\_accuracy: 0.8125  
 Epoch 9/60  
 5/5 [=====] - 8s 1s/step - loss: 0.3527 - accuracy:  
 0.8125 - val\_loss: 0.2903 - val\_accuracy: 0.9062  
 Epoch 10/60  
 5/5 [=====] - 8s 1s/step - loss: 0.3291 - accuracy:  
 0.8625 - val\_loss: 0.2573 - val\_accuracy: 0.8750  
 Epoch 11/60  
 5/5 [=====] - 8s 940ms/step - loss: 0.2518 - accuracy:  
 0.9375 - val\_loss: 0.1856 - val\_accuracy: 0.9375  
 Epoch 12/60  
 5/5 [=====] - 9s 1s/step - loss: 0.2275 - accuracy:  
 0.9187 - val\_loss: 0.1235 - val\_accuracy: 0.9375  
 Epoch 13/60  
 5/5 [=====] - 8s 1s/step - loss: 0.1761 - accuracy:  
 0.9125 - val\_loss: 0.1292 - val\_accuracy: 0.9062  
 Epoch 14/60  
 5/5 [=====] - 6s 828ms/step - loss: 0.1685 - accuracy:  
 0.9500 - val\_loss: 0.1338 - val\_accuracy: 0.9375  
 Epoch 15/60  
 5/5 [=====] - 7s 1s/step - loss: 0.1241 - accuracy:  
 0.9375 - val\_loss: 0.1961 - val\_accuracy: 0.8750  
 Epoch 16/60  
 5/5 [=====] - 8s 1s/step - loss: 0.1392 - accuracy:  
 0.9312 - val\_loss: 0.0911 - val\_accuracy: 0.9688  
 Epoch 17/60  
 5/5 [=====] - 7s 849ms/step - loss: 0.1219 - accuracy:  
 0.9438 - val\_loss: 0.0902 - val\_accuracy: 0.9688  
 Epoch 18/60  
 5/5 [=====] - 8s 1s/step - loss: 0.0904 - accuracy:  
 0.9750 - val\_loss: 0.0655 - val\_accuracy: 1.0000  
 Epoch 19/60  
 5/5 [=====] - 7s 842ms/step - loss: 0.0767 - accuracy:  
 0.9688 - val\_loss: 0.0686 - val\_accuracy: 0.9688  
 Epoch 20/60

5/5 [=====] - 7s 1s/step - loss: 0.0718 - accuracy:  
0.9750 - val\_loss: 0.1184 - val\_accuracy: 0.9062  
Epoch 21/60  
5/5 [=====] - 6s 857ms/step - loss: 0.0676 - accuracy:  
0.9688 - val\_loss: 0.2638 - val\_accuracy: 0.8125  
Epoch 22/60  
5/5 [=====] - 8s 953ms/step - loss: 0.0964 - accuracy:  
0.9750 - val\_loss: 0.1467 - val\_accuracy: 1.0000  
Epoch 23/60  
5/5 [=====] - 9s 1s/step - loss: 0.1204 - accuracy:  
0.9750 - val\_loss: 0.1550 - val\_accuracy: 0.9062  
Epoch 24/60  
5/5 [=====] - 6s 826ms/step - loss: 0.0857 - accuracy:  
0.9812 - val\_loss: 0.0392 - val\_accuracy: 1.0000  
Epoch 25/60  
5/5 [=====] - 7s 1s/step - loss: 0.0617 - accuracy:  
0.9750 - val\_loss: 0.0569 - val\_accuracy: 1.0000  
Epoch 26/60  
5/5 [=====] - 6s 837ms/step - loss: 0.0482 - accuracy:  
0.9937 - val\_loss: 0.0135 - val\_accuracy: 1.0000  
Epoch 27/60  
5/5 [=====] - 7s 998ms/step - loss: 0.0488 - accuracy:  
0.9937 - val\_loss: 0.0490 - val\_accuracy: 1.0000  
Epoch 28/60  
5/5 [=====] - 8s 1s/step - loss: 0.0237 - accuracy:  
1.0000 - val\_loss: 0.0221 - val\_accuracy: 1.0000  
Epoch 29/60  
5/5 [=====] - 6s 821ms/step - loss: 0.0355 - accuracy:  
0.9875 - val\_loss: 0.0199 - val\_accuracy: 1.0000  
Epoch 30/60  
5/5 [=====] - 7s 1s/step - loss: 0.0308 - accuracy:  
0.9937 - val\_loss: 0.0065 - val\_accuracy: 1.0000  
Epoch 31/60  
5/5 [=====] - 6s 841ms/step - loss: 0.0362 - accuracy:  
0.9937 - val\_loss: 0.0394 - val\_accuracy: 1.0000  
Epoch 32/60  
5/5 [=====] - 7s 999ms/step - loss: 0.0258 - accuracy:  
0.9937 - val\_loss: 0.0161 - val\_accuracy: 1.0000  
Epoch 33/60  
5/5 [=====] - 6s 822ms/step - loss: 0.0373 - accuracy:  
0.9875 - val\_loss: 0.0664 - val\_accuracy: 0.9688  
Epoch 34/60  
5/5 [=====] - 8s 901ms/step - loss: 0.0224 - accuracy:  
1.0000 - val\_loss: 0.0334 - val\_accuracy: 1.0000  
Epoch 35/60  
5/5 [=====] - 6s 826ms/step - loss: 0.0299 - accuracy:  
0.9875 - val\_loss: 0.0194 - val\_accuracy: 1.0000  
Epoch 36/60

5/5 [=====] - 7s 822ms/step - loss: 0.0161 - accuracy:  
0.9937 - val\_loss: 0.0125 - val\_accuracy: 1.0000  
Epoch 37/60  
5/5 [=====] - 7s 864ms/step - loss: 0.0177 - accuracy:  
0.9937 - val\_loss: 0.0236 - val\_accuracy: 1.0000  
Epoch 38/60  
5/5 [=====] - 7s 847ms/step - loss: 0.0078 - accuracy:  
1.0000 - val\_loss: 0.0035 - val\_accuracy: 1.0000  
Epoch 39/60  
5/5 [=====] - 9s 1s/step - loss: 0.0202 - accuracy:  
0.9937 - val\_loss: 0.0060 - val\_accuracy: 1.0000  
Epoch 40/60  
5/5 [=====] - 6s 844ms/step - loss: 0.0143 - accuracy:  
1.0000 - val\_loss: 0.0091 - val\_accuracy: 1.0000  
Epoch 41/60  
5/5 [=====] - 9s 1s/step - loss: 0.0067 - accuracy:  
1.0000 - val\_loss: 0.0013 - val\_accuracy: 1.0000  
Epoch 42/60  
5/5 [=====] - 6s 826ms/step - loss: 0.0074 - accuracy:  
1.0000 - val\_loss: 0.0032 - val\_accuracy: 1.0000  
Epoch 43/60  
5/5 [=====] - 9s 1s/step - loss: 0.0114 - accuracy:  
1.0000 - val\_loss: 0.0129 - val\_accuracy: 1.0000  
Epoch 44/60  
5/5 [=====] - 7s 1s/step - loss: 0.0108 - accuracy:  
1.0000 - val\_loss: 0.0088 - val\_accuracy: 1.0000  
Epoch 45/60  
5/5 [=====] - 6s 830ms/step - loss: 0.0054 - accuracy:  
1.0000 - val\_loss: 0.0044 - val\_accuracy: 1.0000  
Epoch 46/60  
5/5 [=====] - 7s 1s/step - loss: 0.0047 - accuracy:  
1.0000 - val\_loss: 0.0070 - val\_accuracy: 1.0000  
Epoch 47/60  
5/5 [=====] - 6s 817ms/step - loss: 0.0026 - accuracy:  
1.0000 - val\_loss: 0.0057 - val\_accuracy: 1.0000  
Epoch 48/60  
5/5 [=====] - 9s 1s/step - loss: 0.0028 - accuracy:  
1.0000 - val\_loss: 0.0018 - val\_accuracy: 1.0000  
Epoch 49/60  
5/5 [=====] - 8s 1s/step - loss: 0.0018 - accuracy:  
1.0000 - val\_loss: 0.0033 - val\_accuracy: 1.0000  
Epoch 50/60  
5/5 [=====] - 7s 834ms/step - loss: 0.0013 - accuracy:  
1.0000 - val\_loss: 0.0016 - val\_accuracy: 1.0000  
Epoch 51/60  
5/5 [=====] - 8s 1s/step - loss: 0.0017 - accuracy:  
1.0000 - val\_loss: 0.0013 - val\_accuracy: 1.0000  
Epoch 52/60

```

5/5 [=====] - 6s 834ms/step - loss: 0.0013 - accuracy:
1.0000 - val_loss: 0.0020 - val_accuracy: 1.0000
Epoch 53/60
5/5 [=====] - 7s 830ms/step - loss: 0.0011 - accuracy:
1.0000 - val_loss: 0.0017 - val_accuracy: 1.0000
Epoch 54/60
5/5 [=====] - 7s 1s/step - loss: 0.0016 - accuracy:
1.0000 - val_loss: 0.0014 - val_accuracy: 1.0000
Epoch 55/60
5/5 [=====] - 6s 830ms/step - loss: 0.0013 - accuracy:
1.0000 - val_loss: 0.0018 - val_accuracy: 1.0000
Epoch 56/60
5/5 [=====] - 8s 1s/step - loss: 0.0012 - accuracy:
1.0000 - val_loss: 2.4022e-04 - val_accuracy: 1.0000
Epoch 57/60
5/5 [=====] - 7s 835ms/step - loss: 9.6568e-04 -
accuracy: 1.0000 - val_loss: 9.5465e-04 - val_accuracy: 1.0000
Epoch 58/60
5/5 [=====] - 7s 841ms/step - loss: 9.1853e-04 -
accuracy: 1.0000 - val_loss: 0.0012 - val_accuracy: 1.0000
Epoch 59/60
5/5 [=====] - 8s 1s/step - loss: 7.4237e-04 - accuracy:
1.0000 - val_loss: 9.1980e-04 - val_accuracy: 1.0000
Epoch 60/60
5/5 [=====] - 7s 837ms/step - loss: 5.8313e-04 -
accuracy: 1.0000 - val_loss: 9.6101e-04 - val_accuracy: 1.0000

```


## 1.4 Plot Performance

```

[29]: fig = plt.figure()
plt.plot(hist.history['loss'], color='teal', label='loss')
plt.plot(hist.history['val_loss'], color='orange', label='val_loss')
fig.suptitle('Loss', fontsize=20)
plt.legend(loc="upper left")
plt.show()

```


## Loss



```
[30]: fig = plt.figure()
plt.plot(hist.history['accuracy'], color='teal', label='accuracy')
plt.plot(hist.history['val_accuracy'], color='orange', label='val_accuracy')
fig.suptitle('Accuracy', fontsize=20)
plt.legend(loc="upper left")
plt.show()
```



## Accuracy



### 1.5 Evaluate

```
[31]: from tensorflow.keras.metrics import Precision, Recall, BinaryAccuracy
```

```
[32]: pre = Precision()  
      re = Recall()  
      acc = BinaryAccuracy()
```

```
[33]: for batch in test.as_numpy_iterator():  
      X, y = batch  
      yhat = model.predict(X)  
      pre.update_state(y, yhat)  
      re.update_state(y, yhat)  
      acc.update_state(y, yhat)
```

```
[34]: print(pre.result(), re.result(), acc.result())
```

```
tf.Tensor(0.0, shape=(), dtype=float32) tf.Tensor(0.0, shape=(), dtype=float32)  
tf.Tensor(0.0, shape=(), dtype=float32)
```

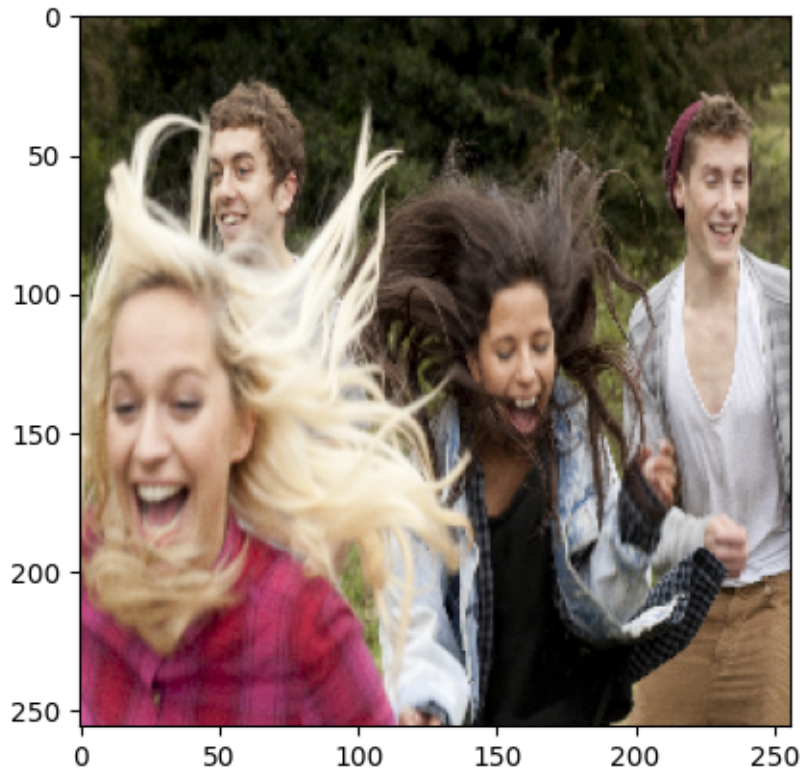
## 1.6 10. Test

```
[35]: import cv2
```

```
[36]: img = cv2.imread('drive/MyDrive/Project 13 Classification/154006829.jpg')
img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
if img is None:
    print("Image not loaded")
else:
    plt.imshow(img)
    plt.show()
```



```
[37]: resize = tf.image.resize(img, (256,256))
plt.imshow(resize.numpy().astype(int))
plt.show()
```



```
[38]: yhat = model.predict(np.expand_dims(resize/255, 0))
```

```
1/1 [=====] - 0s 323ms/step
```

```
[39]: yhat
```

```
[39]: array([[7.7254776e-07]], dtype=float32)
```

```
[40]: if yhat > 0.5:
      print(f'Predicted class is Sad')
      else:
      print (f'Predicted Class is happy')
```

Predicted Class is happy

Function for predicting images if happy or sad

```
[41]: def predict_emotion(image_path, model):
      # Load and convert the image
      img = cv2.imread(image_path)
      if img is None:
          print("Image not loaded")
          return
```

```

img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)

# Resize the image
resize = tf.image.resize(img, (256,256))

# Predict the class
yhat = model.predict(np.expand_dims(resize/255, 0))
print(yhat)
if yhat > 0.5:
    print('Predicted class is Sad')
else:
    print('Predicted class is Happy')

# Show the original and resized images
plt.figure(figsize=(10, 5))
plt.subplot(1, 2, 1)
plt.imshow(img)
plt.title('Original Image')
plt.subplot(1, 2, 2)
plt.imshow(resize.numpy().astype(int))
plt.title('Resized Image')
plt.show()

```

```

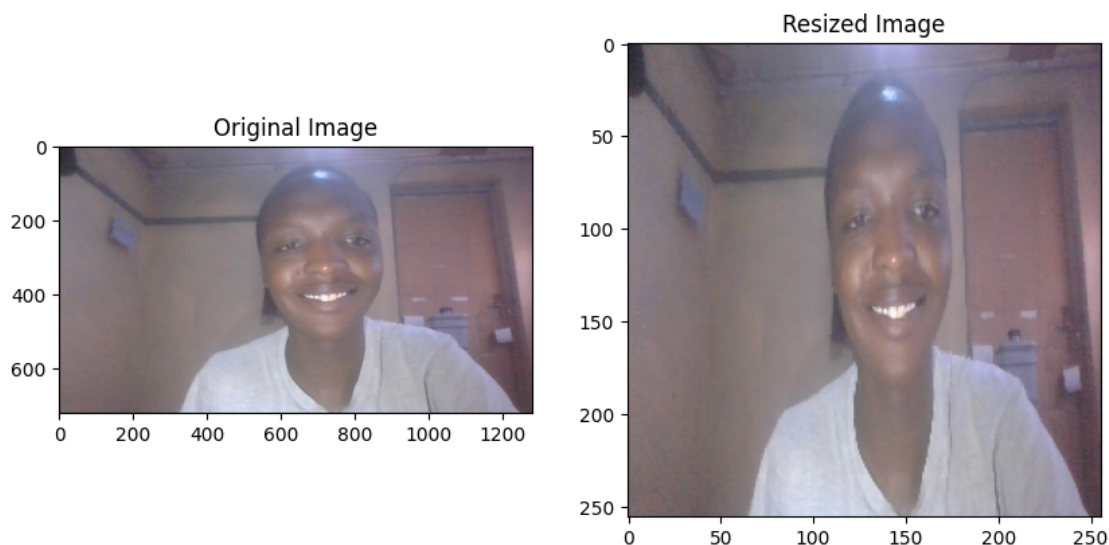
[42]: image_path = 'drive/MyDrive/Project 13 Classification/WIN_20240216_13_46_13_Pro.
      ↪jpg'
      predict_emotion(image_path, model)

```

1/1 [=====] - 0s 36ms/step

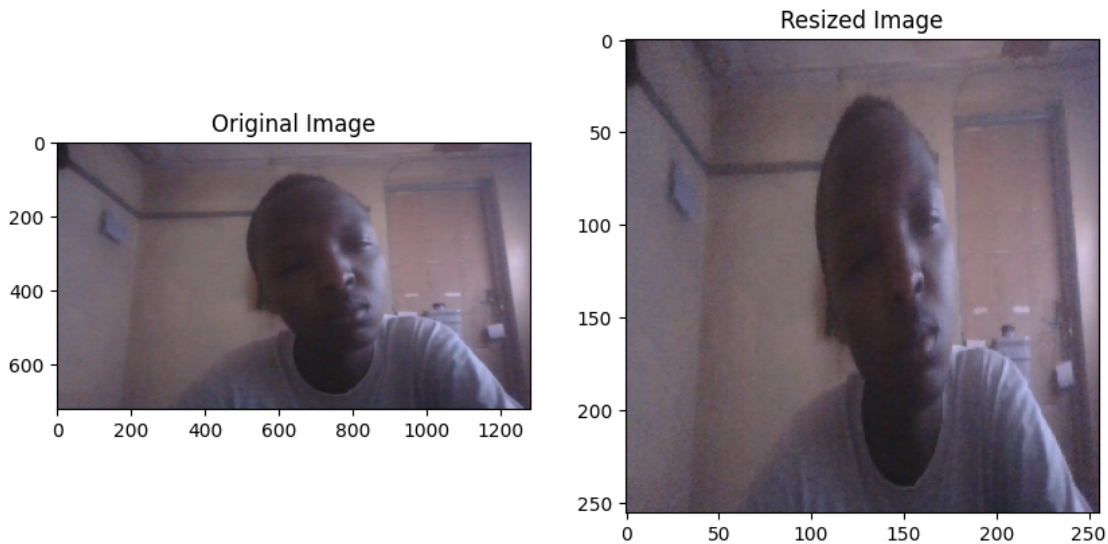
[[0.00483828]]

Predicted class is Happy



```
[43]: image_path = 'drive/MyDrive/Project 13 Classification/WIN_20240216_13_14_07_Pro.
      ↪jpg'
      predict_emotion(image_path, model)
```

```
1/1 [=====] - 0s 18ms/step
[[0.71629643]]
Predicted class is Sad
```



## 1.7 Save The Model

```
[44]: from tensorflow.keras.models import load_model
```

```
[45]: model.save(os.path.join('models', 'drive/MyDrive/Project 13 Classification/
      ↪model/imageclassifier.h5'))
```

```
/usr/local/lib/python3.10/dist-packages/keras/src/engine/training.py:3103:
UserWarning: You are saving your model as an HDF5 file via `model.save()`. This
file format is considered legacy. We recommend using instead the native Keras
format, e.g. `model.save('my_model.keras')`.
  saving_api.save_model(
```